

```
#Storing daily sales data in a NumPy array.
import numpy as np
sales = np.array([150, 200, 250, 300, 400, 350, 500]) # Sales for each day
print("Sales Data:", sales)
```

```
#Calculating mean, maximum, and minimum sales.
```

```
print("Average Sales:", np.mean(sales))
print("Highest Sale:", np.max(sales))
print("Lowest Sale:", np.min(sales))
```

```
#Creating a table-like structure using Pandas.
```

```
import pandas as pd
data = {
    "Customer": ["Alice", "Bob", "Charlie"],
    "Age": [25, 30, 35],
    "Amount Spent": [120, 200, 150]
}
df = pd.DataFrame(data)
print(df)
```

```
⇒ Sales Data: [150 200 250 300 400 350 500]
Average Sales: 307.14285714285717
Highest Sale: 500
Lowest Sale: 150
   Customer  Age  Amount Spent
0    Alice   25           120
1     Bob    30           200
2  Charlie   35           150
```

```
#Uploading and reading a CSV file in Pandas.
```

```
import pandas as pd
from google.colab import files
uploaded = files.upload() # Upload your CSV file
df = pd.read_csv("students_dataset.csv") # Read and loads the file into a Pandas DataFrame.
df.head()
```

```
# Check for missing values
```

```
print(df.isnull().sum())
```

```
#Remove Missing Values
```

```
df_cleaned = df.dropna() # Removes rows with missing values
print(df_cleaned)
```

```
#Fill with Mean/Median (For Numerical Data)
```

```
df["Age"].fillna(df['Age'].mean(),inplace=True)
df["Marks"].fillna(df["Marks"].median(), inplace=True)
df["Attendance"].fillna(df["Attendance"].mean(),inplace=True)
#Fill with Mode (For Categorical Data)
df["Passed"].fillna(df["Passed"].mode()[0], inplace=True)
print(df)
```

```
df.ffill(inplace=True) # Forward fill
```

```
df.bfill(inplace=True) # Backward fill
```

```
#Remove Duplicates
```

```
df.drop_duplicates(inplace=True)
```

```

df.drop_duplicates(inplace=True)

#Standardization (Z-score Normalization)
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
df_scaled=df.copy()
df_scaled[["Marks", "Attendance"]]=scaler.fit_transform(df[["Marks", "Attendance"]])
print(df_scaled)

#Min-Max Scaling
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()

df_scaled[["Marks", "Attendance"]]=scaler.fit_transform(df[["Marks", "Attendance"]])
print(df_scaled)

#One-Hot Encoding
df_encoded = pd.get_dummies(df, columns=["Passed"],drop_first=True)
print(df_encoded)

#Label Encoding
from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
df["Passed"] = encoder.fit_transform(df["Passed"])
print(df)

#Creating New Features

def performance_category(marks):
    if marks>=85:
        return "High"
    elif marks >= 70:
        return "Medium"
    else:
        return "Low"

df["Performance"] = df["Marks"].apply(performance_category)
print(df)
#Binning (Converting Continuous to Categorical Data)

df["Age_Group"] = pd.cut(df["Age"], bins=[18, 21, 24],labels=["Young", "Adult"])
print(df)

```



Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current

browser session. Please rerun this cell to enable.

Saving students_dataset.csv to students_dataset.csv

```
Name      0
Age        4
Marks      4
Attendance 4
Passed     0
```

dtype: int64

	Name	Age	Marks	Attendance	Passed
0	Alice	20.0	85.0	90.0	Yes
3	David	22.0	90.0	85.0	Yes
4	Eve	20.0	88.0	95.0	No
8	Ivy	24.0	92.0	92.0	Yes
12	Mona	21.0	83.0	80.0	No

	Name	Age	Marks	Attendance	Passed
0	Alice	20.0	85.0	90.00	Yes
1	Bob	21.0	82.0	80.00	No
2	Charlie	22.0	78.0	85.25	Yes
3	David	22.0	90.0	85.00	Yes
4	Eve	20.0	88.0	95.00	No
5	Frank	22.0	76.0	70.00	Yes
6	Grace	23.0	82.0	88.00	No
7	Hank	21.0	80.0	85.25	Yes
8	Ivy	24.0	92.0	92.00	Yes
9	Jack	22.0	79.0	85.00	No
10	Kelly	22.0	82.0	78.00	Yes
11	Leo	23.0	87.0	85.25	Yes
12	Mona	21.0	83.0	80.00	No
13	Nina	22.0	77.0	89.00	Yes
14	Oscar	25.0	82.0	91.00	No
15	Paul	22.0	81.0	85.25	Yes

<ipython-input-3-8ba44832badd>:18: FutureWarning: A value is trying to be set on a copy of a D. The behavior will change in pandas 3.0. This inplace method will never work because the interm

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: valu

```
df["Age"].fillna(df['Age'].mean(),inplace=True)
```

<ipython-input-3-8ba44832badd>:19: FutureWarning: A value is trying to be set on a copy of a D. The behavior will change in pandas 3.0. This inplace method will never work because the interm

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: valu

```
df["Marks"].fillna(df["Marks"].median(), inplace=True)
```

<ipython-input-3-8ba44832badd>:20: FutureWarning: A value is trying to be set on a copy of a D. The behavior will change in pandas 3.0. This inplace method will never work because the interm

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: valu

```
df["Attendance"].fillna(df["Attendance"].mean(),inplace=True)
```

<ipython-input-3-8ba44832badd>:22: FutureWarning: A value is trying to be set on a copy of a D. The behavior will change in pandas 3.0. This inplace method will never work because the interm

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: valu

```
import matplotlib.pyplot as plt
```

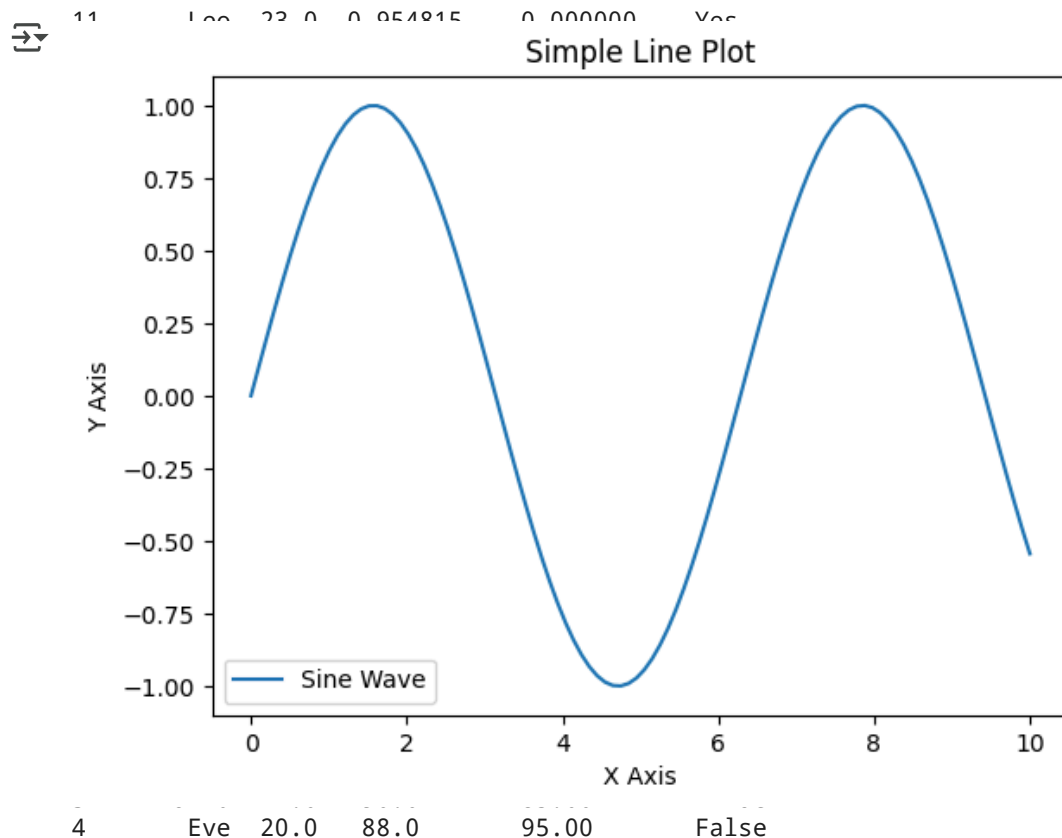
```
import numpy as np
```

```
x = np.linspace(0, 10, 100)
```

```
y = np.sin(x)
```

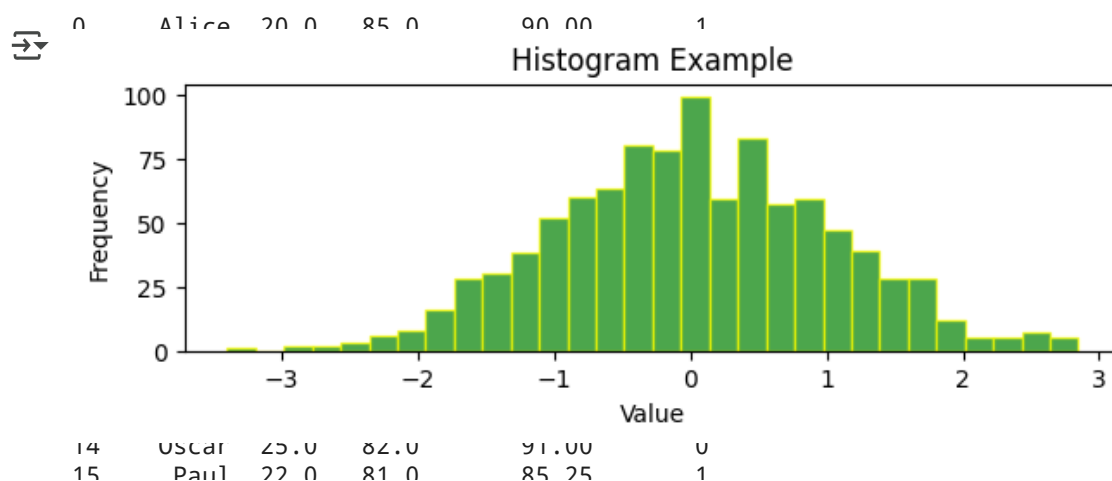
```
plt.plot(x, y, label="Sine Wave")
```

```
plt.xlabel("X Axis")
plt.ylabel("Y Axis")
plt.title("Simple Line Plot")
plt.legend()
plt.show()
```



#1.3 Histogram (Distribution of Data)

```
data = np.random.randn(1000)
#print(data)
plt.figure(figsize=(7, 2))
plt.hist(data, bins=30, color='green', edgecolor='yellow', alpha=0.7)
plt.xlabel("Value")
plt.ylabel("Frequency")
plt.title("Histogram Example")
plt.show()
```




#2. Seaborn - Statistical Data Visualization ,2.1 Histogram & KDE Plot

```
import seaborn as sns
import pandas as pd
```

```
# Creating sample data
data = np.random.randn(1000)
df = pd.DataFrame(data, columns=['Values'])
print(df['Values'])
# Plot
sns.histplot(df['Values'], bins=10, kde=True, color='blue')
plt.title("Histogram with KDE")
plt.show()
```

```
14 Oscar 25.0 82.0 91.00 0 Medium
15 Paul 22.0 81.0 85.25 1 Medium
16 Nam 25.0 82.0 85.00 1 High
17 Alice 20.0 85.0 90.00 1 High
18 Bob 21.0 82.0 80.00 0 Medium
19 Charlie 22.0 78.0 85.25 1 Medium
20 David 22.0 90.0 85.00 1 High
21 Eve 20.0 88.0 95.00 0 High
22 Frank 20.0 76.0 70.00 1 Medium
23 Grace 20.0 82.0 88.00 0 Medium
24 Hank 21.0 80.0 85.25 1 Medium
25 Ivy 20.0 92.0 92.00 1 High
Name: Values, Length: 1000, dtype: float64
```



#2.2 Box Plot (Detecting Outliers)

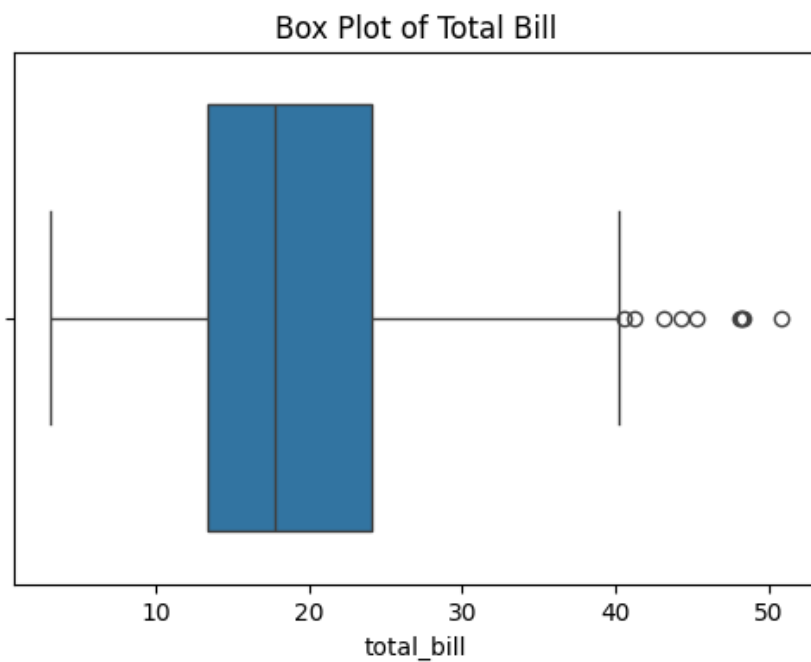
```
tips = sns.load_dataset('tips')
print(tips)
plt.figure(figsize=(6, 4))
sns.boxplot(x=tips['total_bill'])
plt.title("Box Plot of Total Bill")
plt.show()
```

```

⇒
   total_bill  tip    sex smoker  day    time  size
0      16.99  1.01  Female    No   Sun  Dinner    2
1      10.34  1.66    Male    No   Sun  Dinner    3
2      21.01  3.50    Male    No   Sun  Dinner    3
3      23.68  3.31    Male    No   Sun  Dinner    2
4      24.59  3.61  Female    No   Sun  Dinner    4
..      ...    ...    ...    ...   ...    ...    ...
239     29.03  5.92    Male    No   Sat  Dinner    3
240     27.18  2.00  Female   Yes   Sat  Dinner    2
241     22.67  2.00    Male   Yes   Sat  Dinner    2
242     17.82  1.75    Male    No   Sat  Dinner    2
243     18.78  3.00  Female    No  Thur  Dinner    2

```

[244 rows x 7 columns]

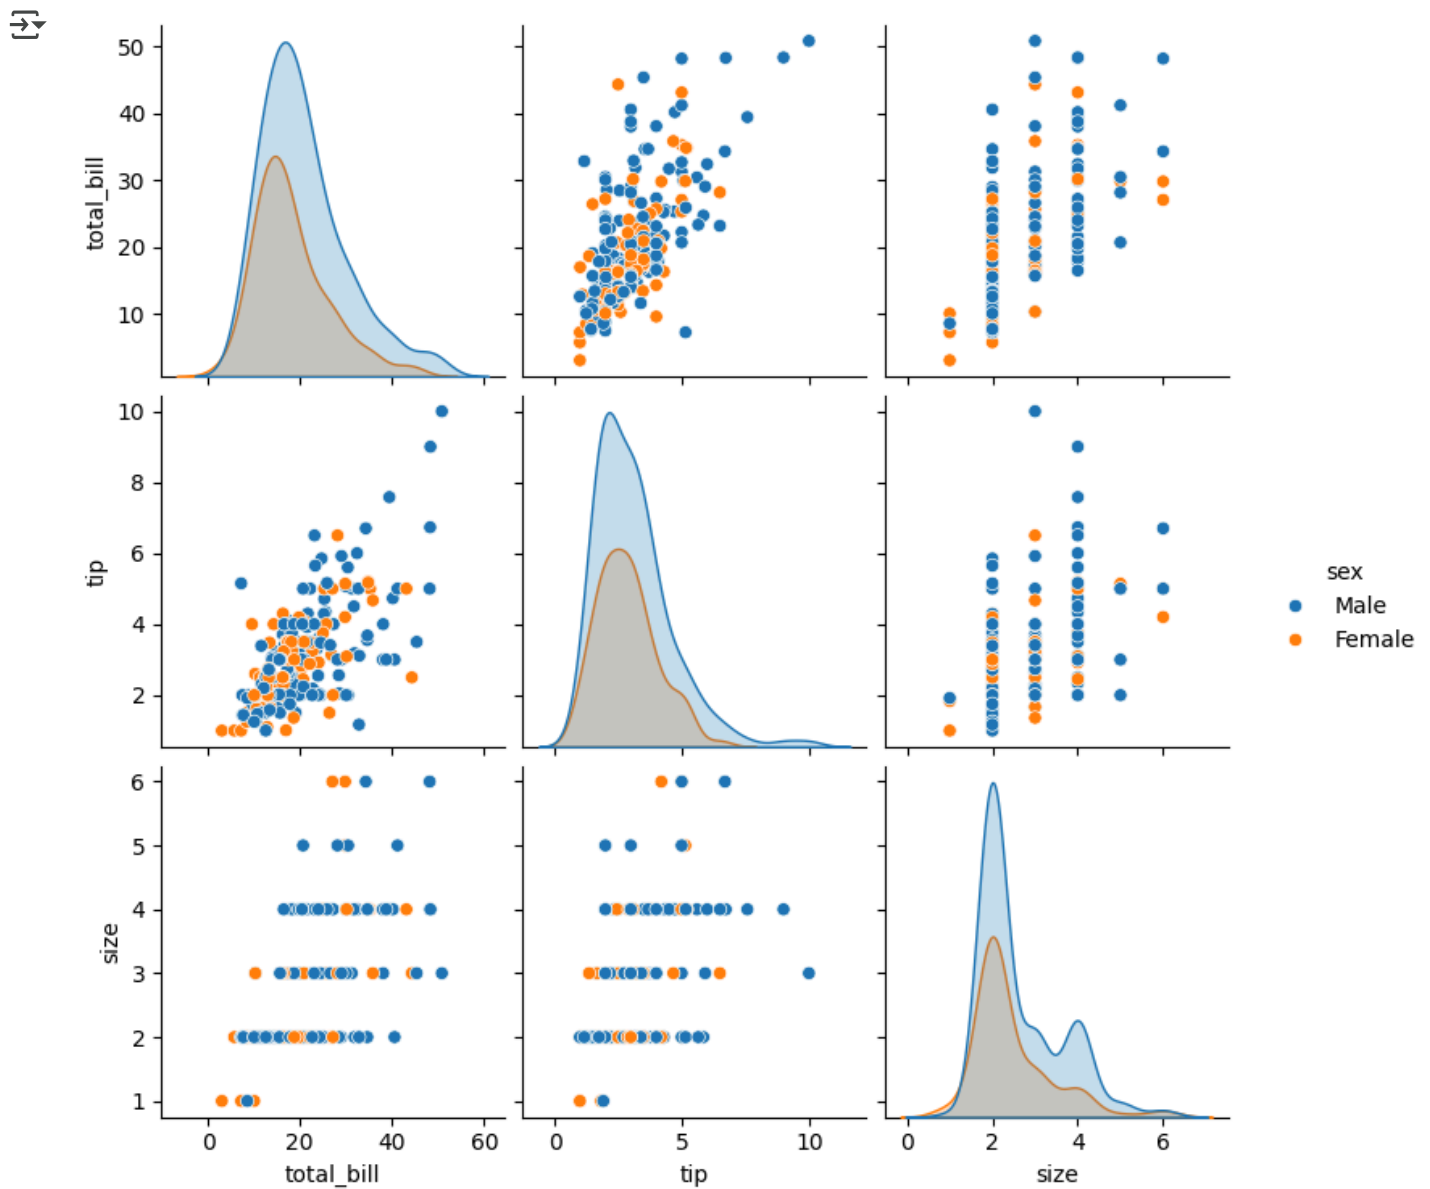


#2.3 Pair Plot (Exploring Relationships)

```

sns.pairplot(tips, hue='sex')
plt.show()

```




#2.4 Heatmap (Correlation Analysis)

```
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
tips = sns.load_dataset('tips')

# Select only numeric columns for correlation
numeric_tips = tips.select_dtypes(include='number')
corr_matrix = numeric_tips.corr()

# Plot the heatmap
plt.figure(figsize=(7, 5))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



#3. Pyplot,3.1 Interactive Line Plot


```
import pandas as pd
import numpy as np
import plotly.express as px

df = pd.DataFrame({"x": np.linspace(0, 10, 100), "y": np.sin(np.linspace(0, 10, 100))
})

fig = px.line(df, x='x', y='y', title="Interactive Sine Wave")
fig.show()
```




Interactive Sine Wave



#3.3 Interactive 3D Scatter Plot

```
import plotly.graph_objects as go
```

```
fig = go.Figure(data=[go.Scatter3d(
    x=tips['total_bill'],
    y=tips['tip'],
    z=tips['size'],
    mode='markers',
    marker=dict(size=5, color=tips['total_bill'], colorscale='Viridis')
)])
```

```
fig.update_layout(title="3D Scatter Plot of Total Bill, Tip & Size")
fig.show()
```